

50 Shades of Concurrency

About me

About me

@michaelklishin

About me



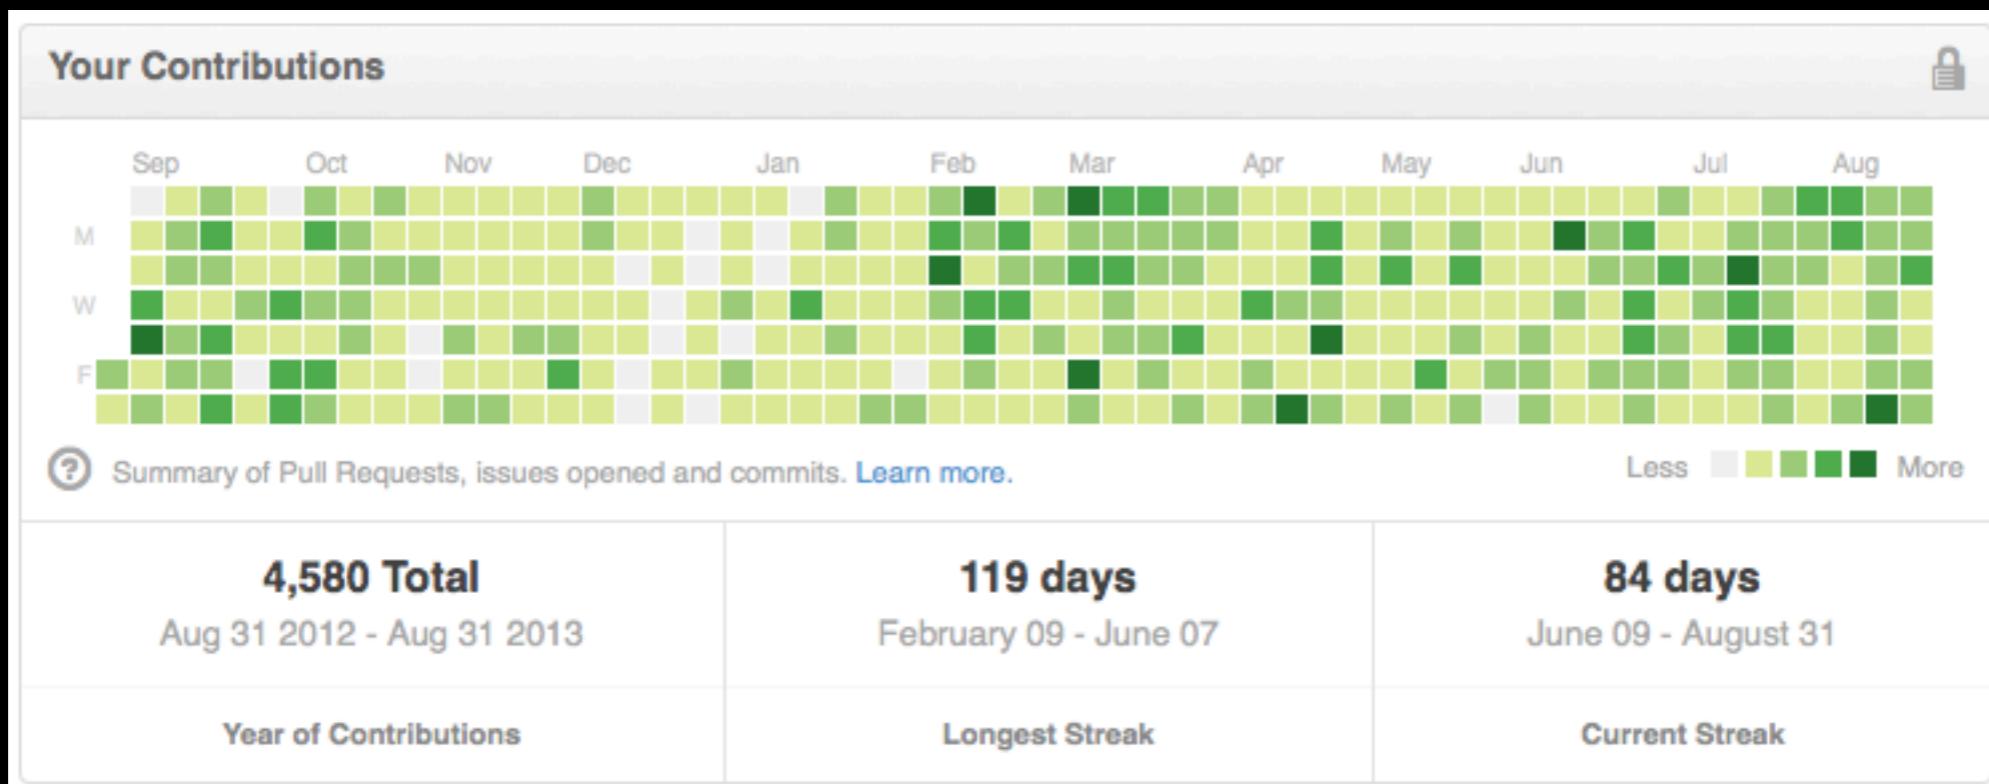
About me

clojurewerkz.org

About me

Ex-Travis CI
core

About me



About me

data processing
and analytics

About me

dozen of programming
languages, <3 λ

About me

recognized animated
gifs expert

About this talk

About this talk

N lightning talks
combined

About this talk

Multiple related
topics

About this talk

Observations

About this talk

Rants

Do we really need concurrency?



The shit is just too
damn hard

Thank you

but no



Do we need concurrency?



«In 2013 alone, IP traffic is expected to grow by around 14 exabytes per month...» — ITU

Exabyte? LOLWUT?

Exabyte

- 10^{18} bytes

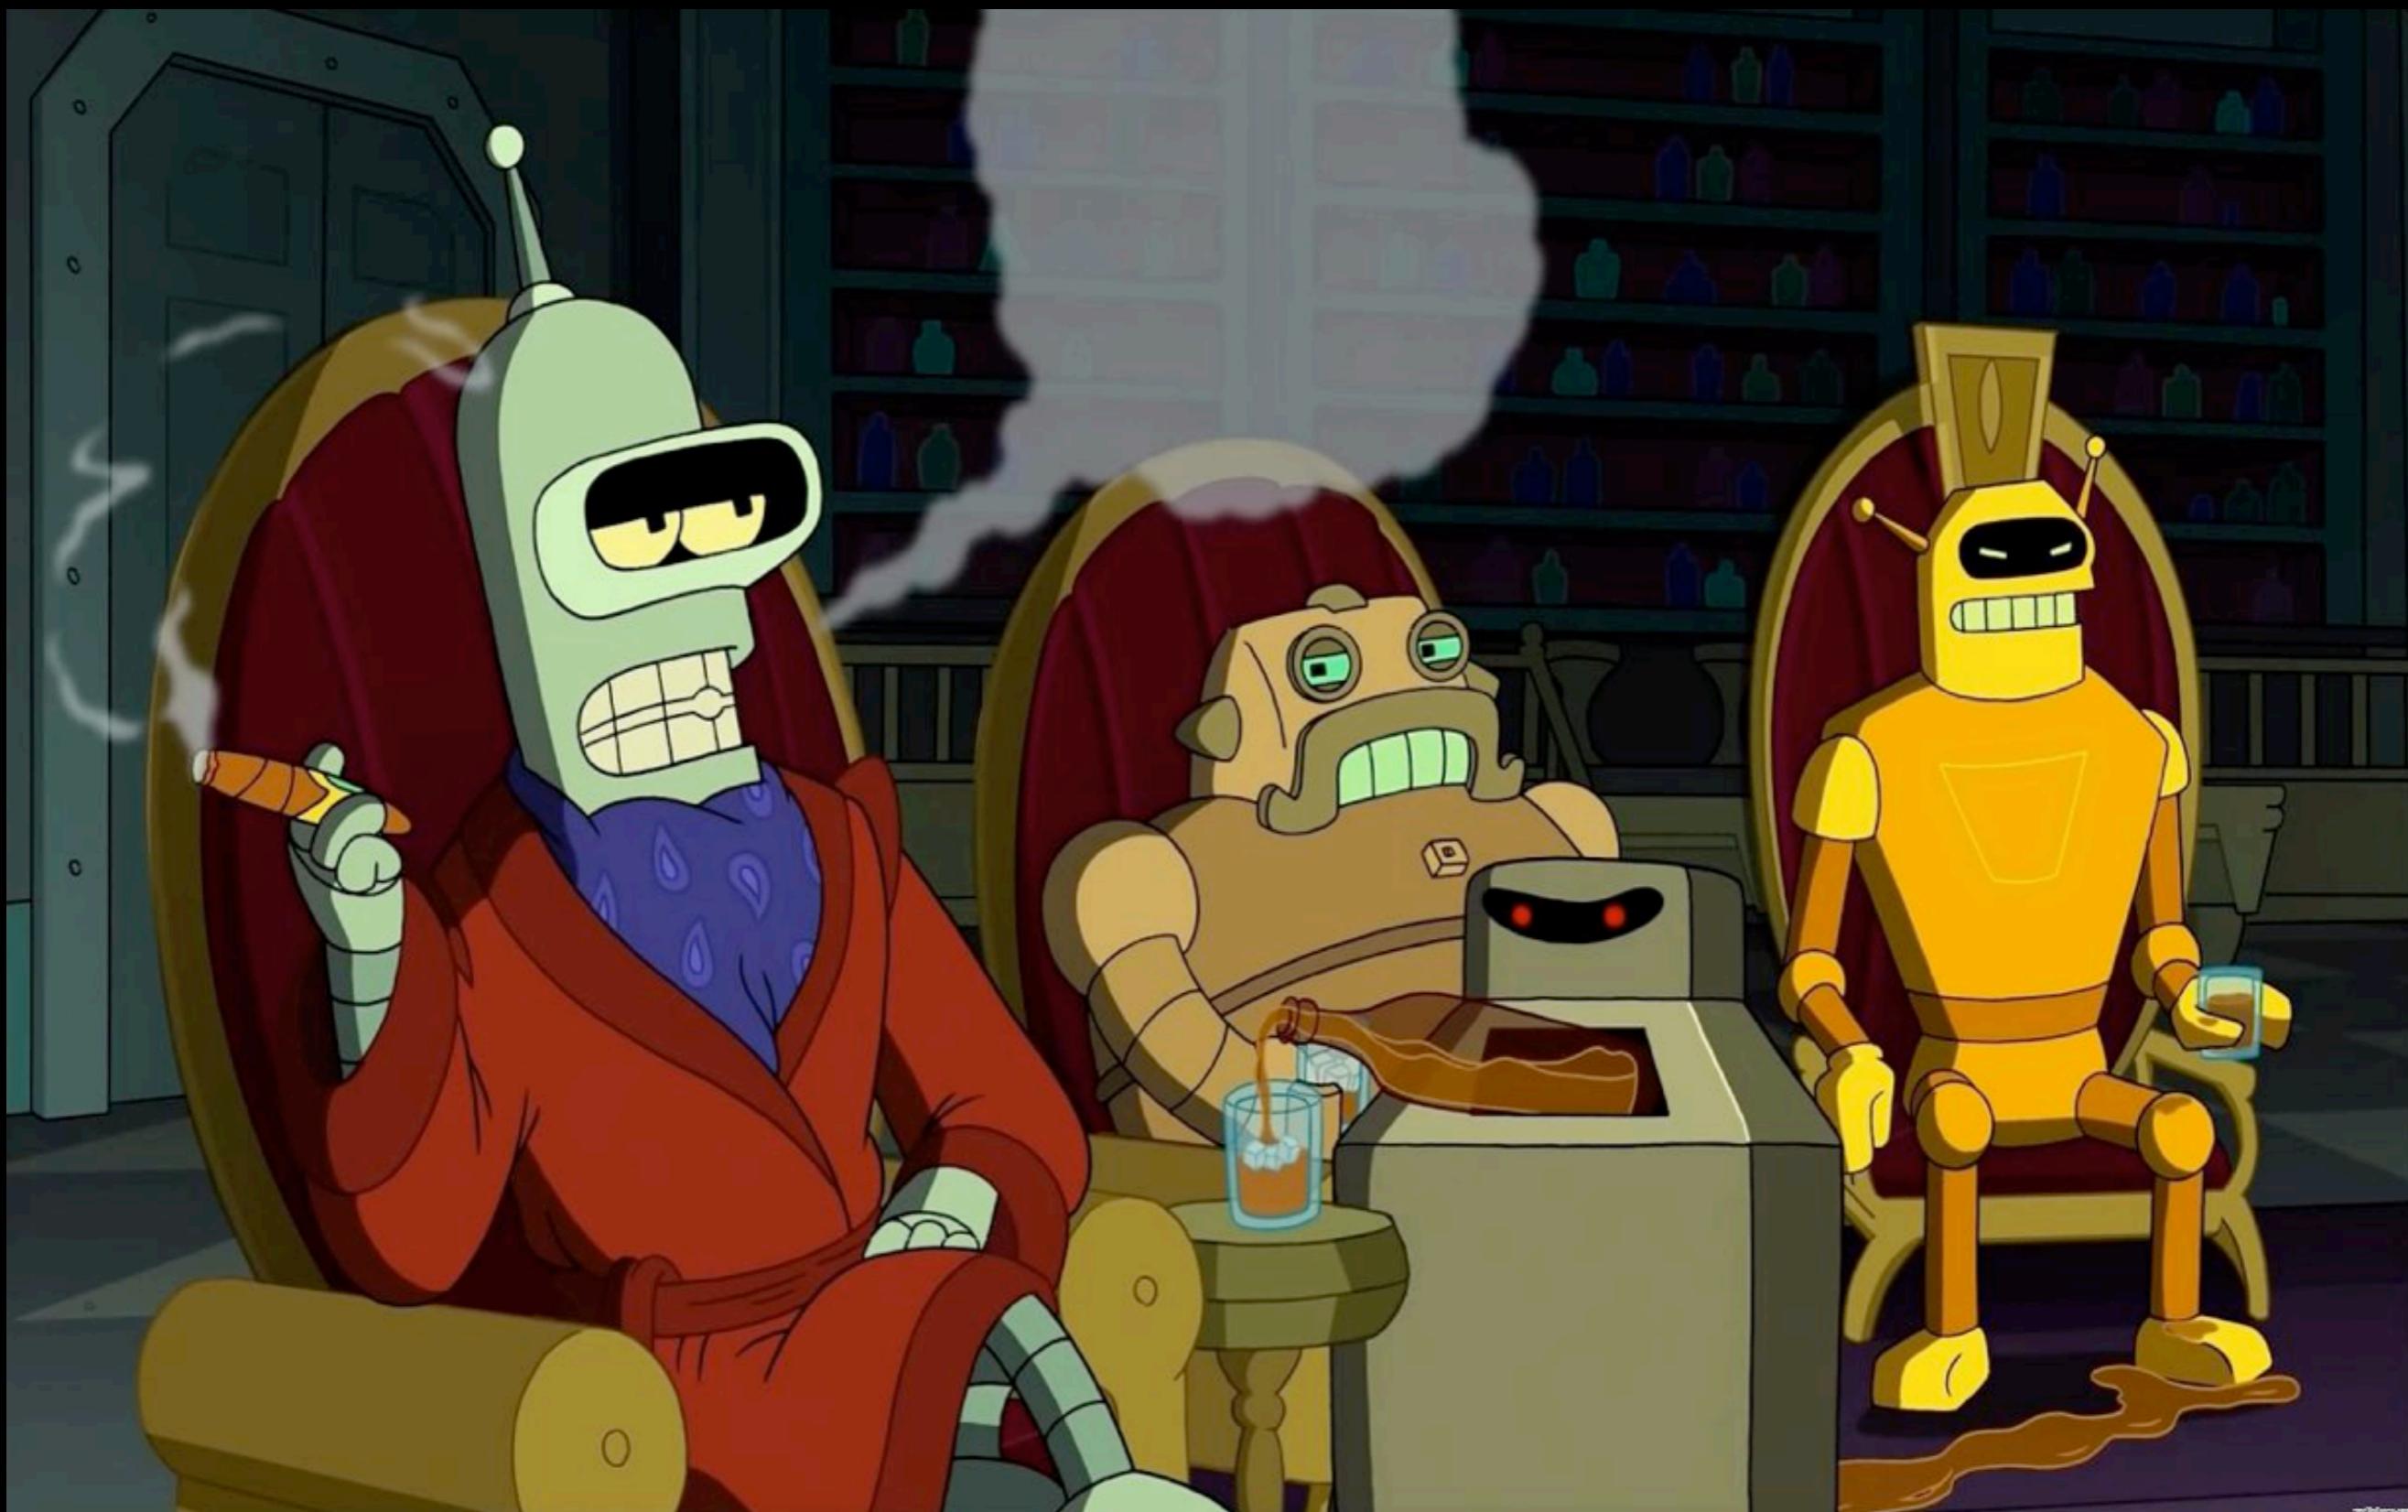
Exabyte

- 10^{18} bytes
- 1,000 petabytes

Exabyte

- 10^{18} bytes
- 1,000 petabytes
- 1,000,000,000 gigabytes





If your project is data-driven, concurrency is no longer a nice-to-have

Do we need concurrency?

Michael Klishin

[Repositories](#) [Profile](#)

Flick the switches below to turn on the Travis service hook for your projects, then push to GitHub.

Last synchronized from GitHub: about 16 hours ago

[Sync now](#)

michaelklishin/acits

OFF

michaelklishin/adventures_with_ssl_talk

OFF

michaelklishin/amqp_broker_stress_tests

OFF

michaelklishin/apache-jackrabbit-chef-cookbook

OFF

michaelklishin/cassandra-chef-cookbook

OFF

michaelklishin/chash

OFF

michaelklishin/clang-chef-cookbook

OFF

michaelklishin/clj1062

OFF

michaelklishin/clojure-cookbook

OFF

michaelklishin/community

OFF

michaelklishin/crawlista

ON

michaelklishin/cucumber.el

OFF

michaelklishin/cyclist

ON

Timely updates from N sources
is a UI feature

Products compete on
user experience

Reasonably responsive
UI or **GTFO**

Do we need concurrency?

The world around us is
concurrent

Pull iPhone headphones
jack when the track is
being switched

Why it happens?

Race condition

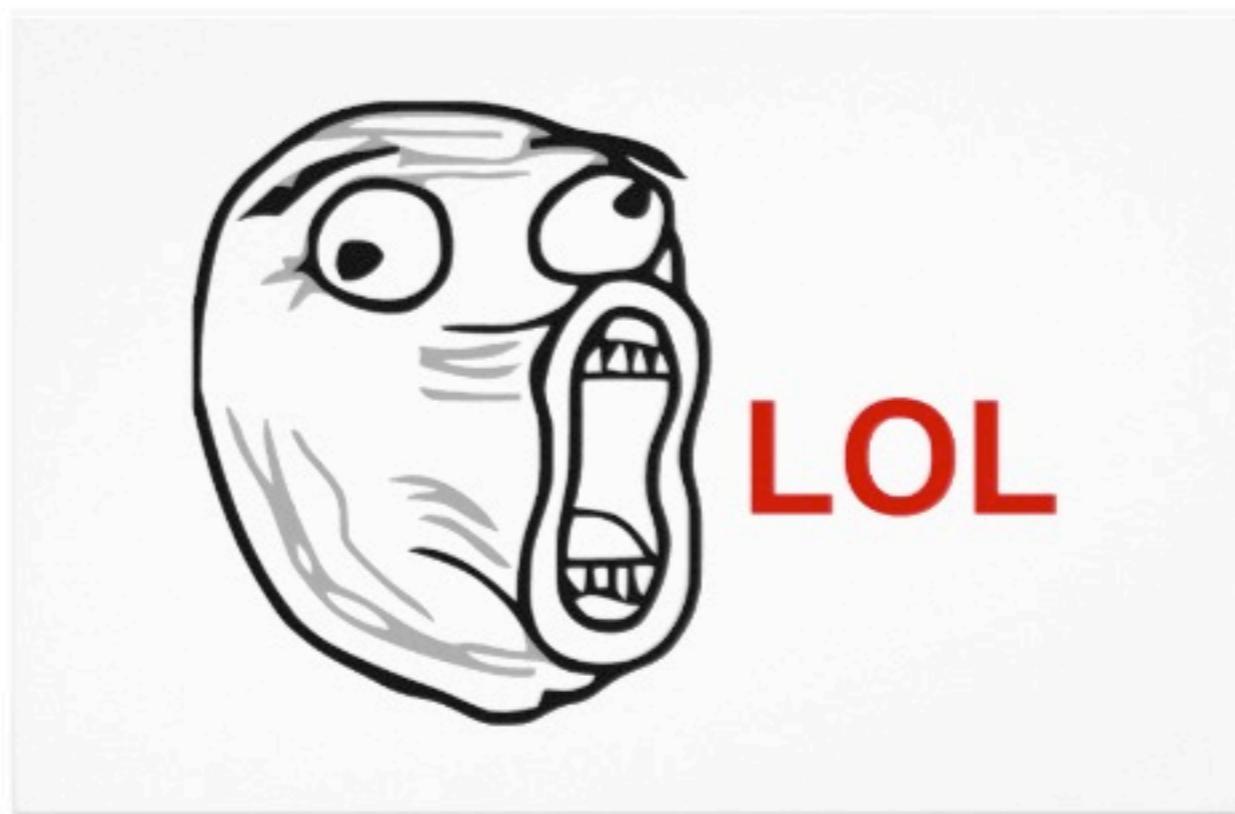
Race conditions, explained



Do we need concurrency?

Do we need to
understand
concurrency?

«my framework will
solve this problem for
me!»



```
require "mutex"
```

```
m = Mutex.new
```

```
m.synchronize do
```

```
  m.synchronize do
```

```
    # ...
```

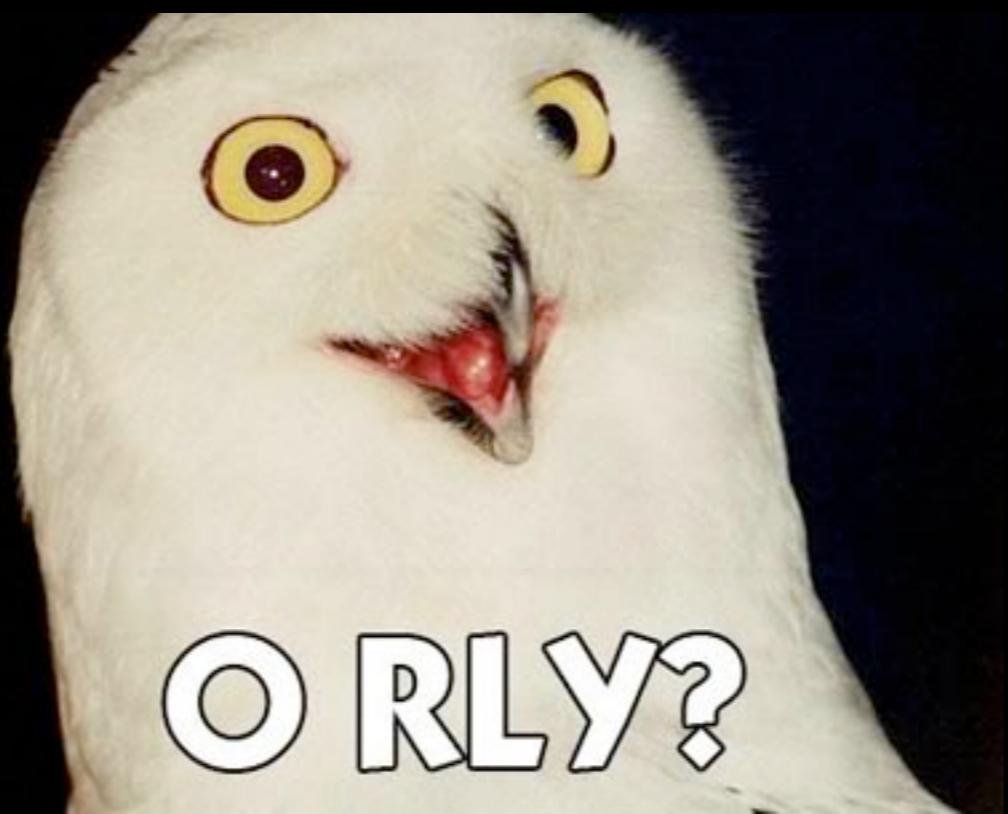
```
end
```

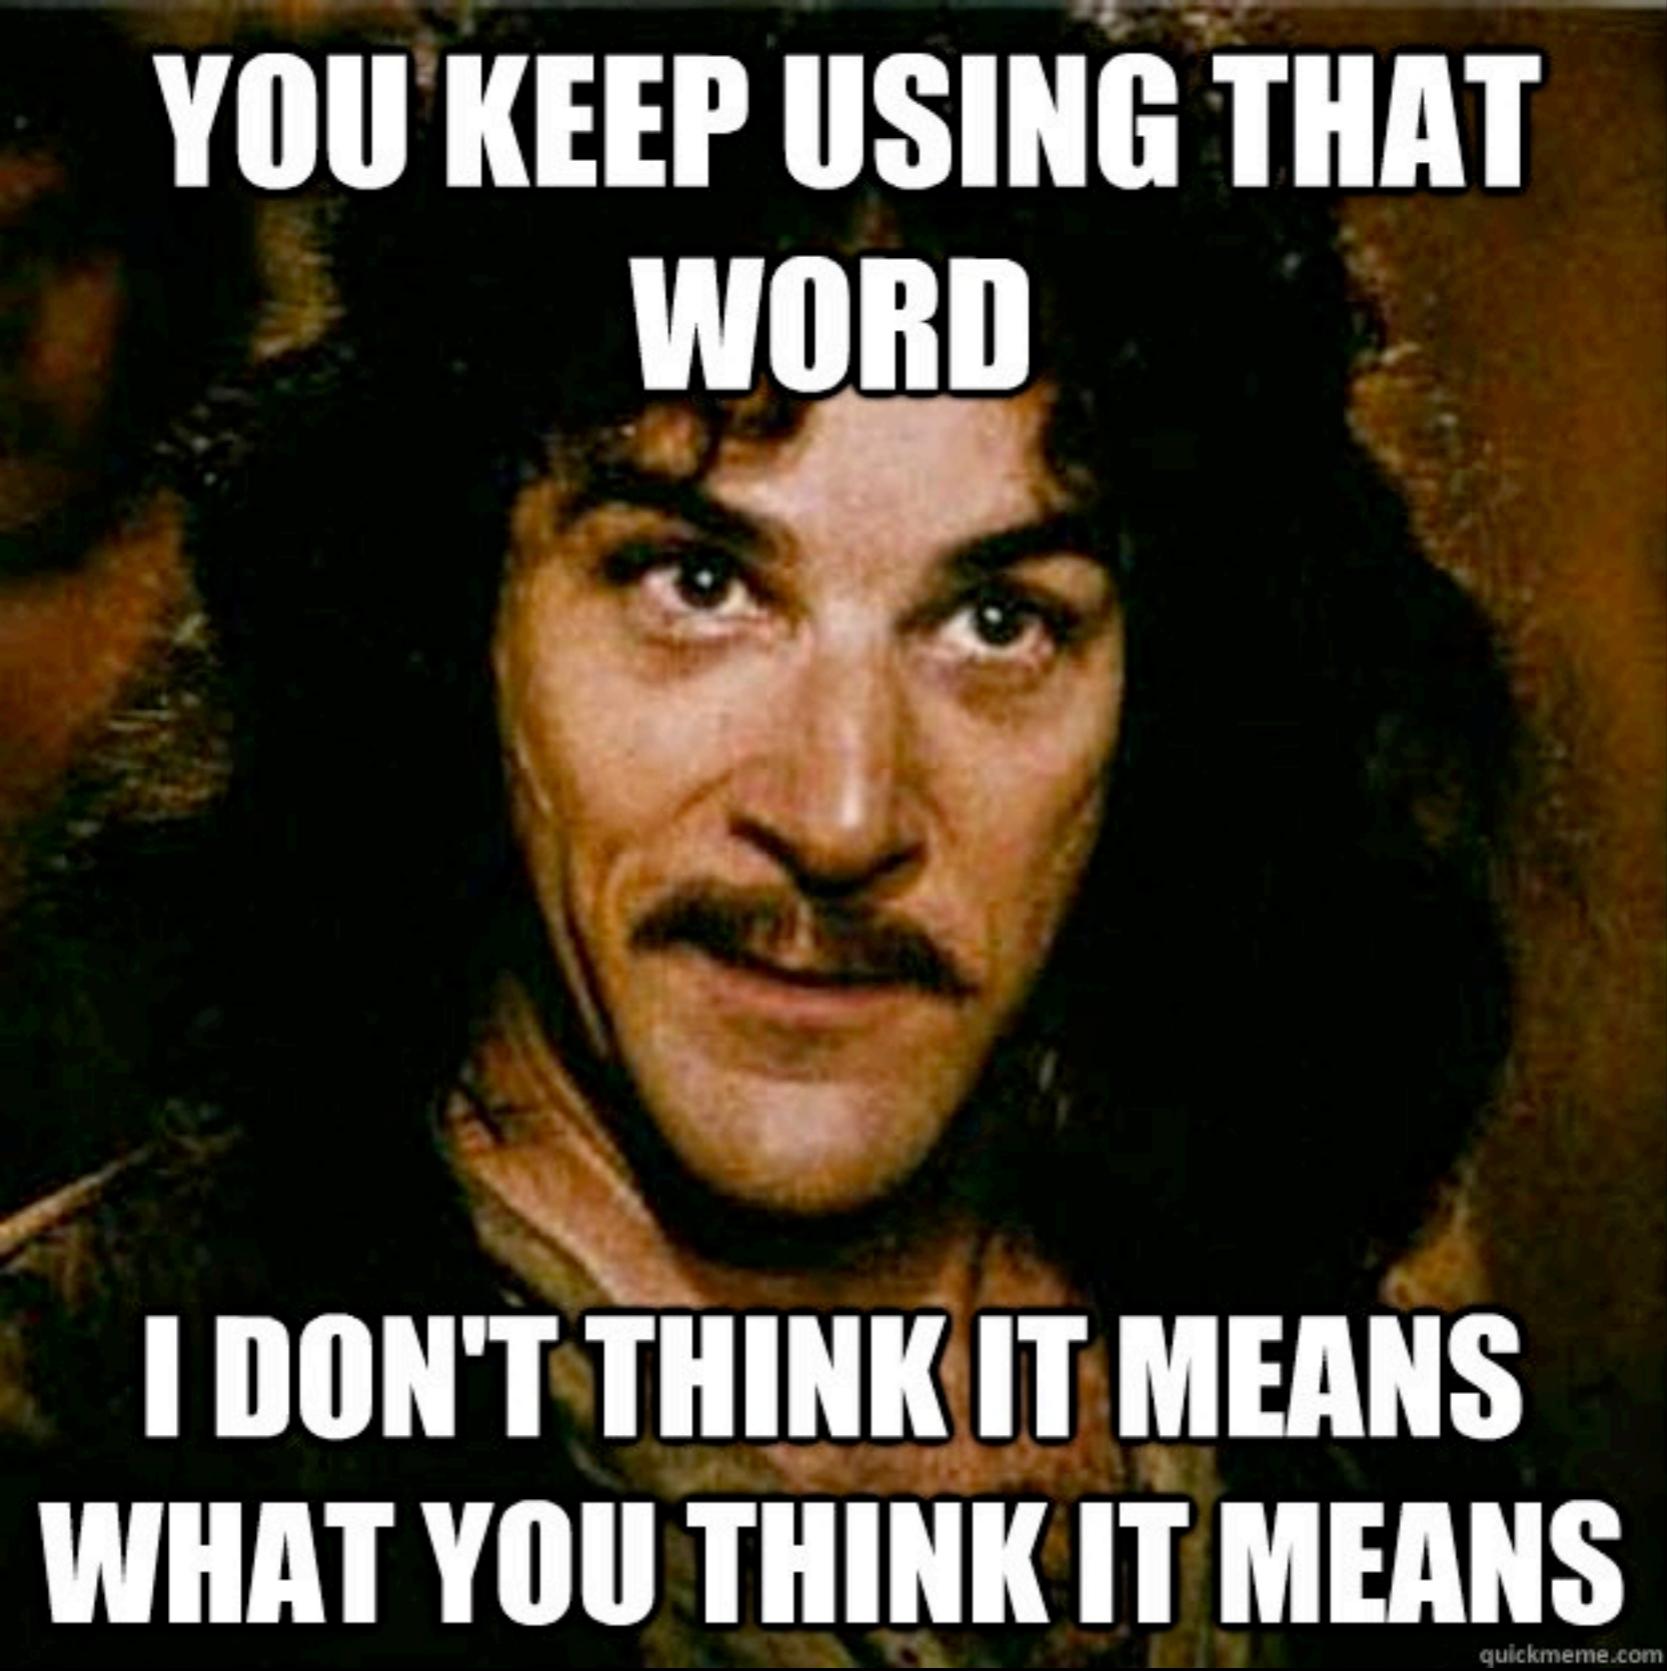
```
end
```



ThreadError: **deadlock; recursive locking**

```
from (irb):4:in `synchronize'  
from (irb):4:in `block in  
irb_binding'  
from (irb):4:in `synchronize'
```





**YOU KEEP USING THAT
WORD**

**I DON'T THINK IT MEANS
WHAT YOU THINK IT MEANS**

quickmeme.com

«Deadlock describes a situation when one or more threads are blocked **waiting on** each other...»

```
require "mutex"
```

```
m = Mutex.new
```

```
m.synchronize do
```

```
  m.synchronize do
```

```
    # ...
```

```
end
```

```
end
```

Unfortunately, Ruby standard library has more **poor design decisions** w.r.t. concurrency

Solvable?

require “monitor”

5-10% throughput
decrease

Backwards compatibility



Mutexes, explained

Mutexes, explained



«I know, Java is
incredibly solid and
stable!»

```
Map<String, Integer> m =  
new HashMap<>();  
  
m.put("michael", 3);
```

results in . . .

CPU usage spikes

HashMap can **create infinite loops internally** if you write to it concurrently, and a write causes a resize.

When I find myself in times of
trouble

Prof. Doug Lea comes to me

Coding lines of wisdom

j.u.C.

java.util.concurrent

«My VM will solve the
problem for me!»

Very GC/allocations
heavy algorithm
implemented in Clojure

«abyssmal multicore
performance, especially
on AMD processors...»

<https://groups.google.com/d/topic/clojure/48W2eff3caU/discussion>

~100 replies and 20
participants later...

The bottleneck is...

OpenJDK/Oracle JDK memory allocator

LOLWUT?

«It's C++ and a few billions of \$
worth of engineering man-hours!»

Not every
algorithm has
concurrent
alternatives

Efficient, concurrent
writes is a hard
problem

Lock contention

Lock contention, explained

Lock contention, explained



Onlylolgifs.tumblr.com

Memory models

«My compiler will solve
the problem for me!»

Compilers are very sophisticated tools and can do a lot for you.
But **they still need your help** from time to time.

False sharing





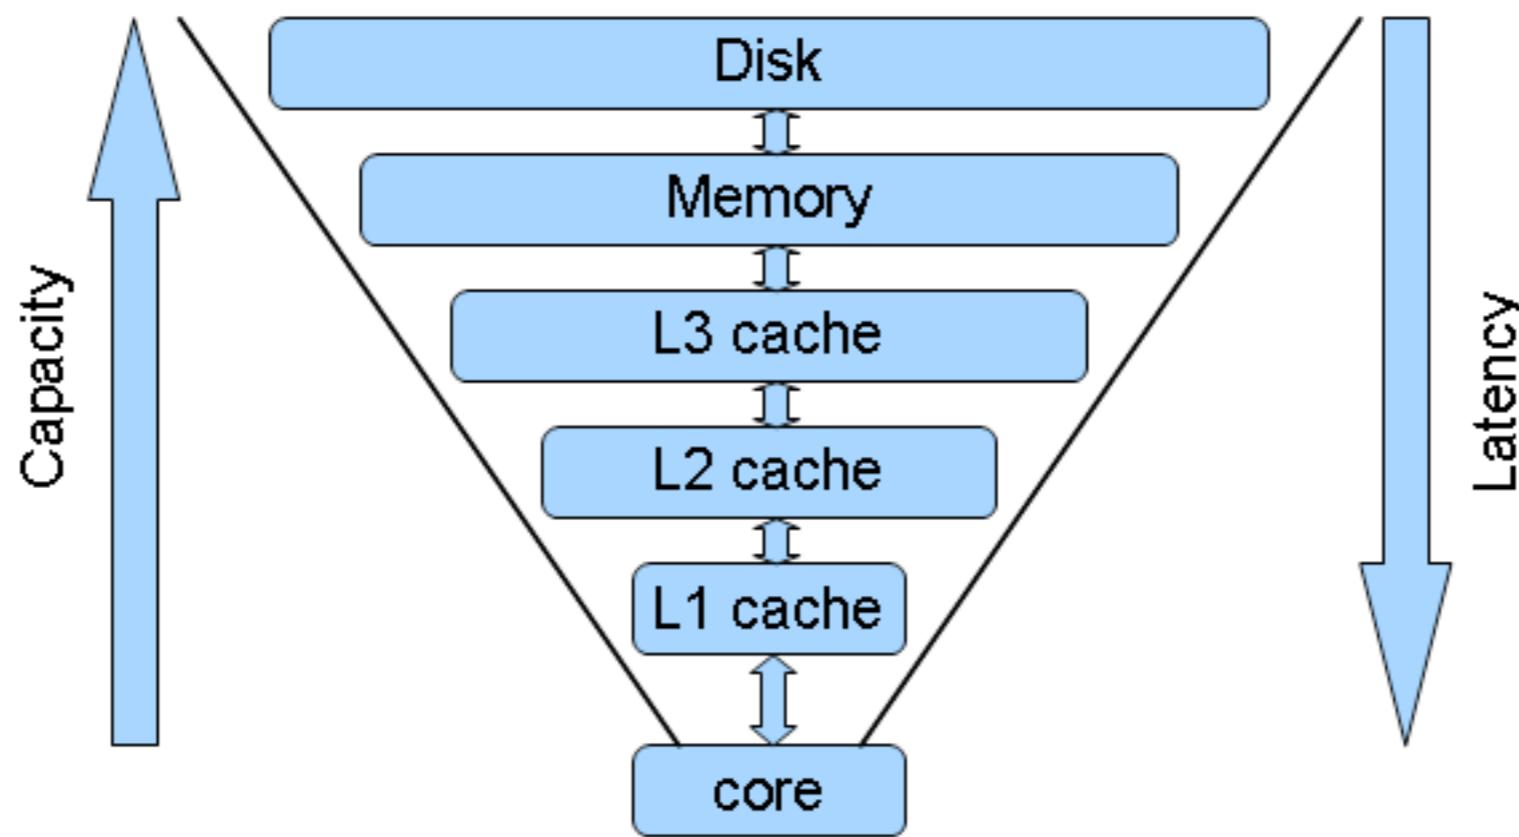
False sharing

```
class Metrics
  attr_accessor :signin_requests
  attr_accessor :signout_requests
end
```

Thread A updates :signin_requests

Thread B updates :signout_requests

CPU caches

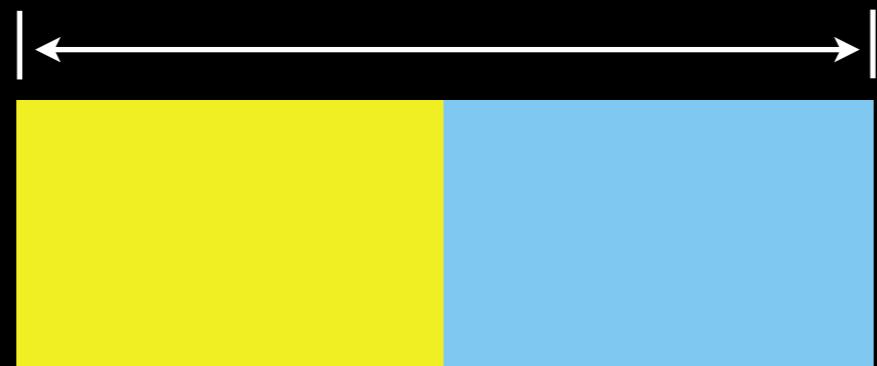


Memory is stored in
CPU caches in **chunks**

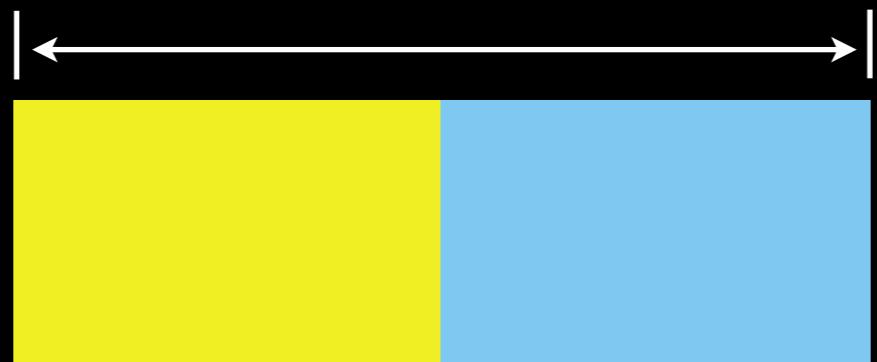
The chunks are known
as cache lines



Cache line

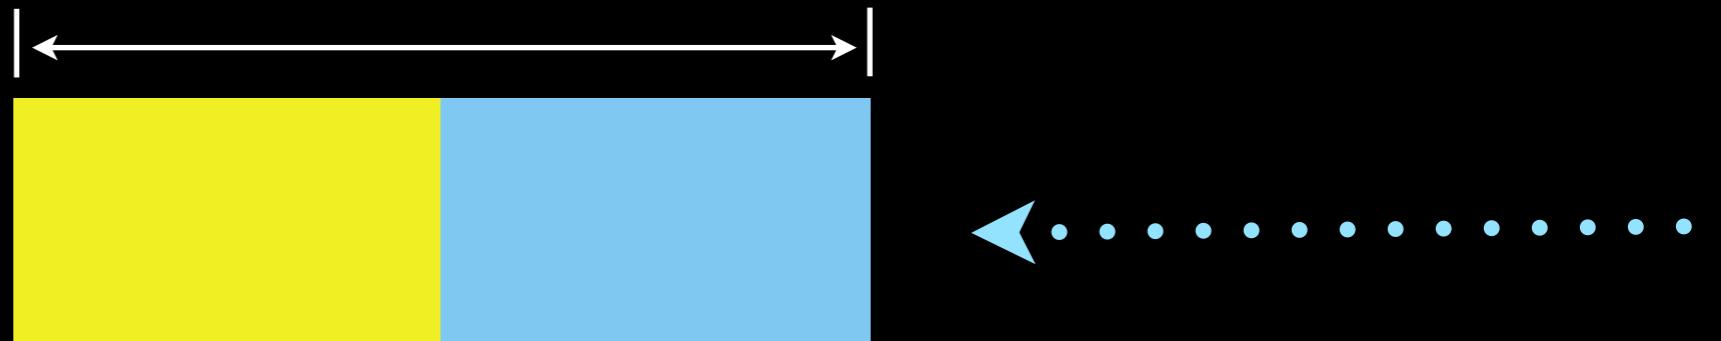


Cache line



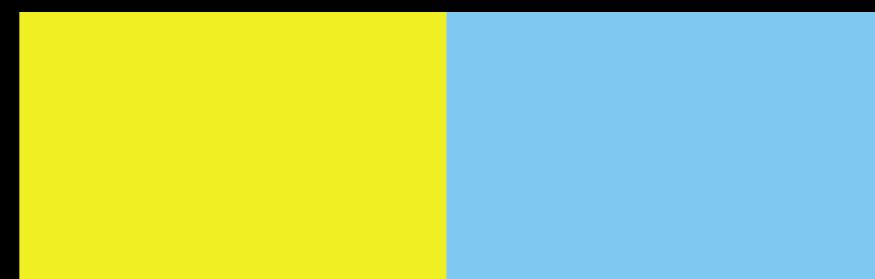
Thread A
updates

Cache line



Thread B
updates

RAM

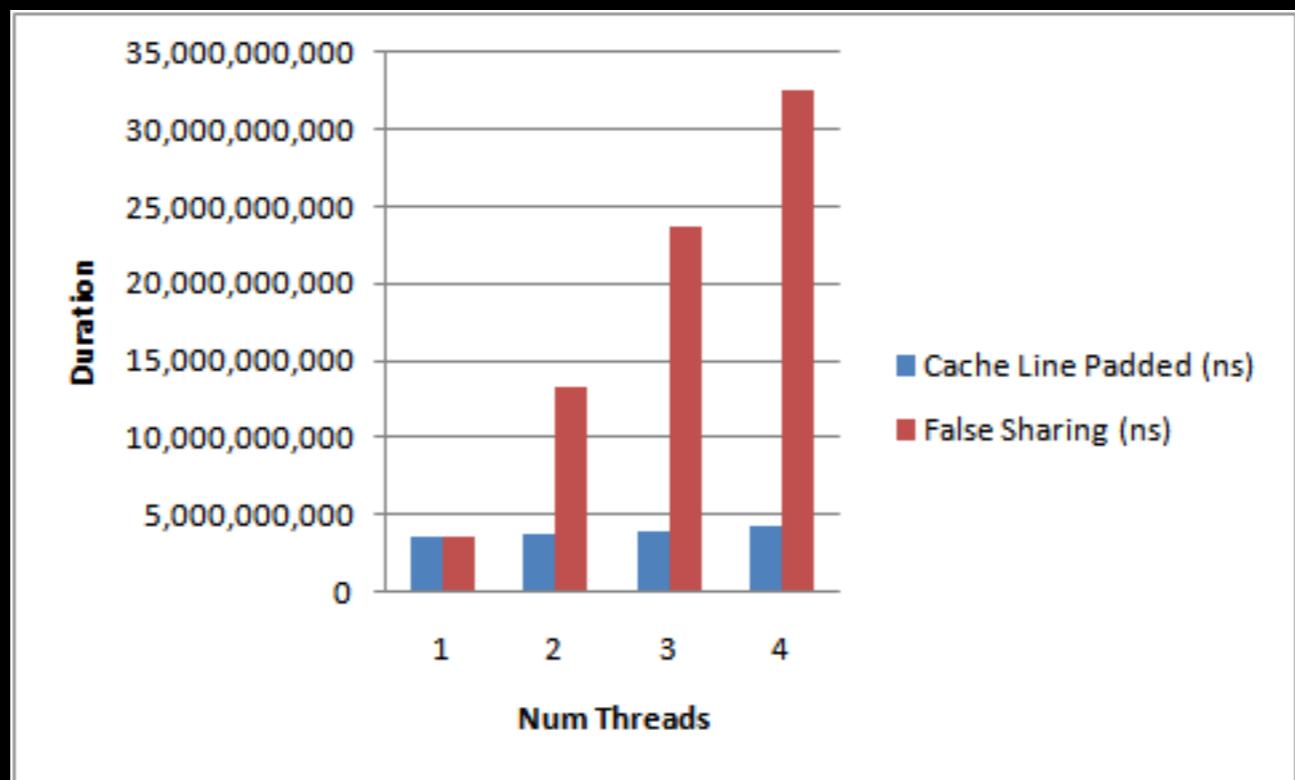


Cache line

Threads update **independent variables**,
but the entire cache line is
has to be read from RAM

Solvable?

```
public final class VolatileLong {  
    public volatile long value = 0L;  
    public long p1, p2, p3, p4, p5, p6;  
}
```



«Threads suck! UNIX processes will
solve my problem!»

fork(2)

fork(2) vs forkall(2)

**fork(2) and parent process
that holds mutex(es)**

async signal safe

«The child inherits copies of the parent's set of open file descriptors...»

«All APIs, including global data symbols, in any framework or library **should be assumed to be unsafe after a fork()** unless explicitly documented to be safe or async-signal safe...»

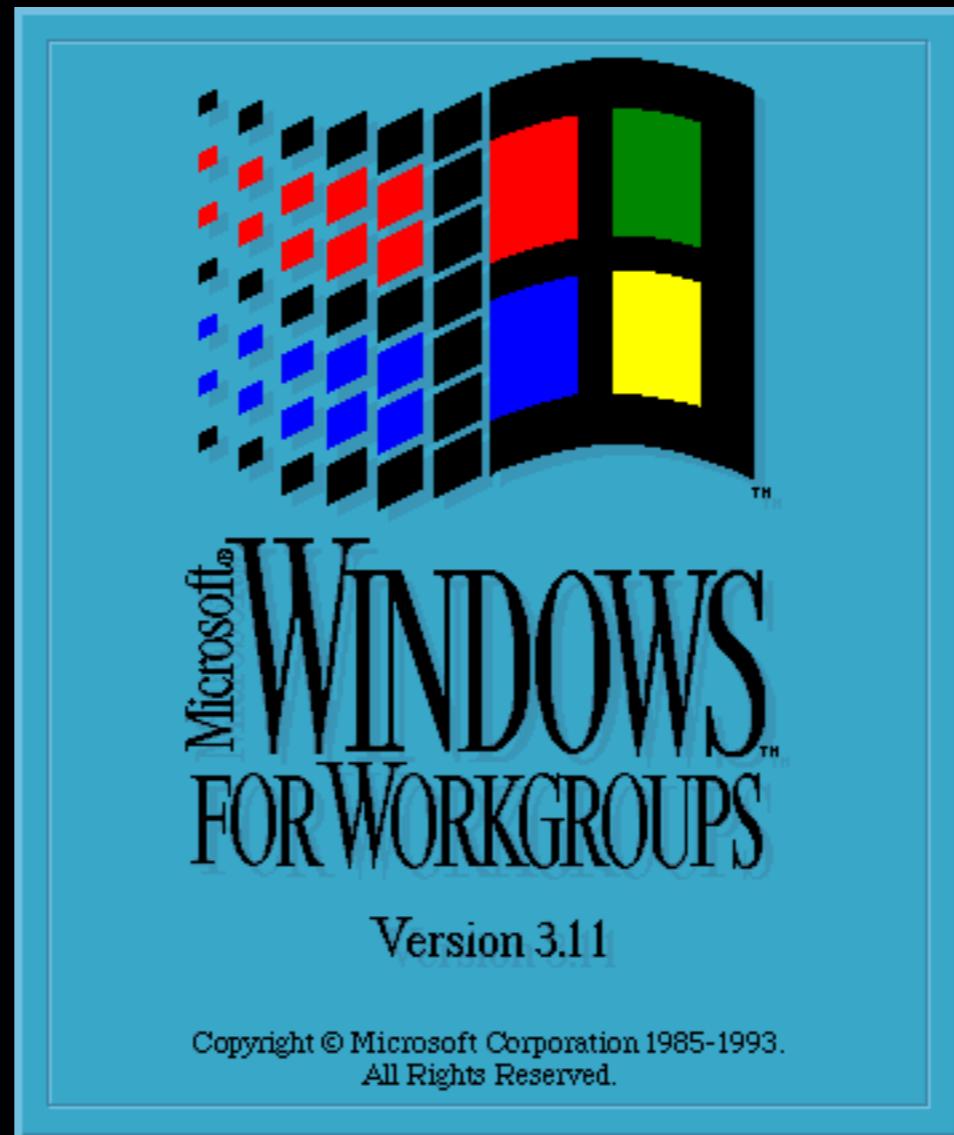
man 2 fork, CAVEATS

man 2 vfork

Process race conditions

Not that many concurrency
libraries for processes

«My OS will solve the
problem for me!»





Cooperative concurrency, explained

Cooperative concurrency, explained



Preemptive scheduling

Scheduling

CIOK

C100M

Lightweight processes

Erlang, Clojure's core.async, Go

Lightweight processes: the Holy Grail?

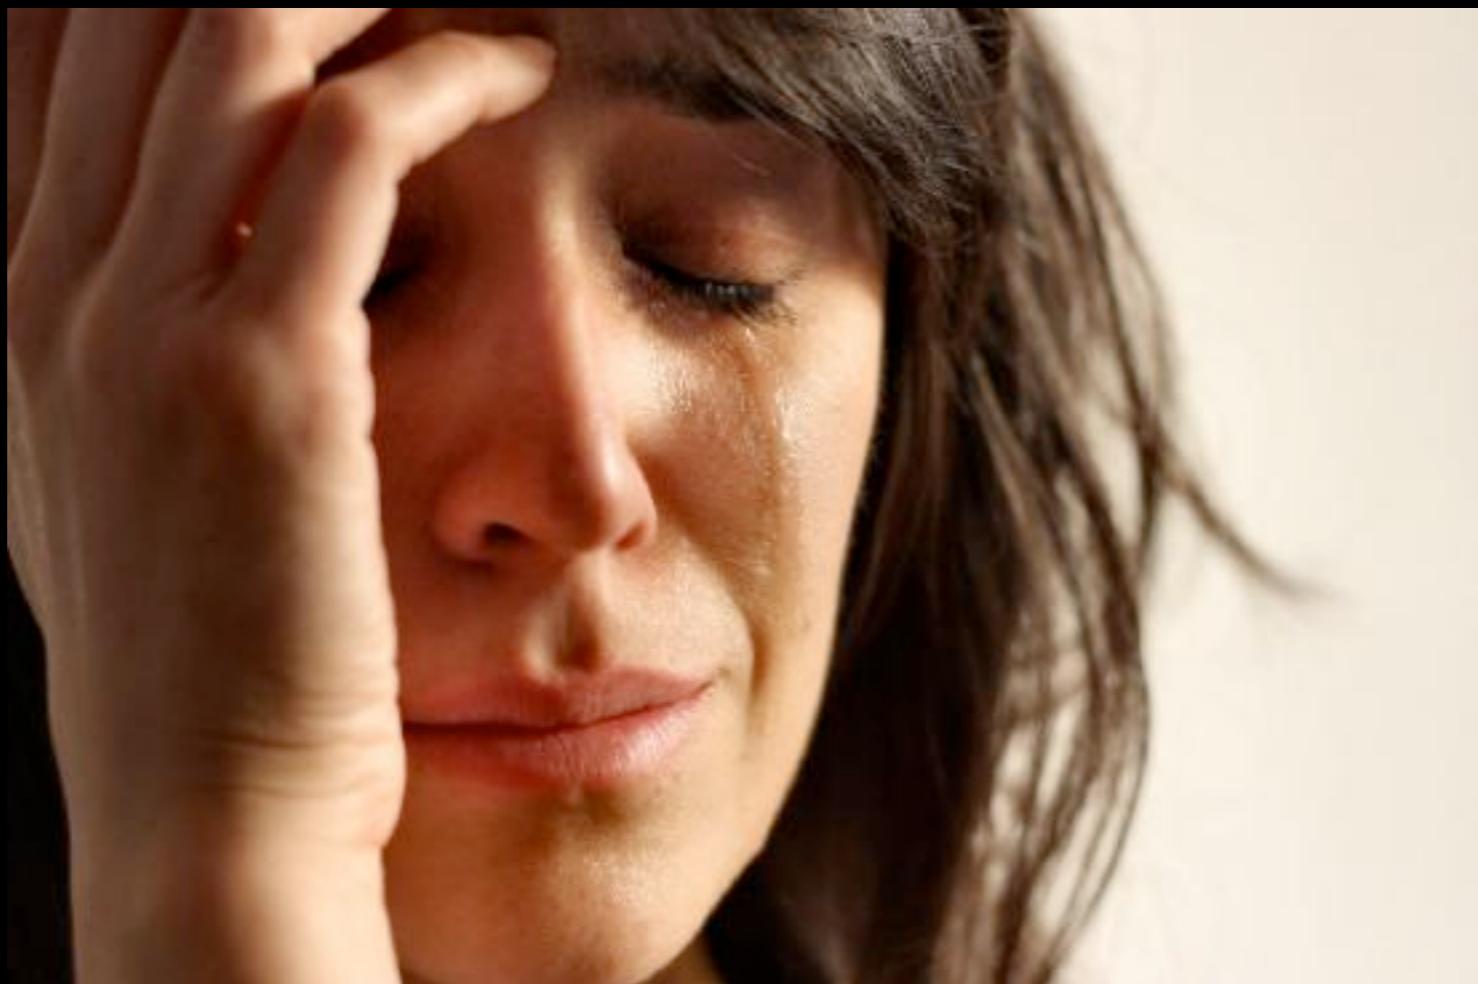
Dat single Erlang
process bottleneck

Unbounded mailboxes and backpressure

Race conditions

Deadlocks in Erlang

Livelocks in Erlang



Concurrency is about
algorithms

Correctness

Race conditions are
possible with any
language

Concurrency makes it
harder to reason about
correctness

Most issues arise from
shared mutable state

Solvable?

Solutions

- Eliminate shared state

Solutions

- Eliminate shared state
- Use immutable data structures

«Immutable? But how
do I change anything?»

Immutable data
structures **create**
copies when updated
(mutated)

«But copying is
expensive!»

Copying doesn't have
to be a naïve “full copy”

Value immutability vs reference immutability

Value immutability: Haskell, Clojure, Erlang.

Reference immutability: C++, Java, Scala, Rust.

Value immutability:
libraries for any
language (Ruby, Java, JS)

Immutability and
message passing
make reasoning
much easier

Summary

Concurrency is
no longer a
nice-to-have

There is no such
thing as a “thread
safe framework”

Concurrency is
about
algorithms

Algorithms are
different at
every level

Problems are
different at
every level

No substitute
for
understanding

Being able to reason
about your program
matters

Immutability and
message passing
make reasoning
much easier

Not about event
loops or Web
Scale™

Concurrency is the new
Cloud Computing

Concurrency is the new
Clown Computering

«Blah blah, X blows Y out of the
water in concurrency...»

Be specific or GTFO

Start learning a
language with
immutable data
structures today

closure-doc.org

Thank you

@michaelklishin

clojurewerkz.org